

ML Library

WOC 6.0


Prepared by: Harshvardhan Saini
(23JE0398)

Linear Regression

I shuffled the given dataset and split it into train, cross validation and test sets in a 0.8 : 0.1 : 0.1 proportion in order to check for overfitting (high variance problem) using the cross validation set and having a neutral check on performance using the test set; carried out shuffling in order to avoid biases in a set. Then, checked whether the dataset is distributed normally or not, in order to check for outliers. I used Z-score normalization to normalize each feature so as to ensure smooth descent in gradient.

Then, i tried different values for the hyperparameter learning rate and ran gradient descent for different number of iterations, and obtained great results for learning rate=1, the cost (Mean Squared Error) reached the global minima: 0.005078585402287006 within 20 iterations.

Cost vs Iterations:



Initial Cost: 65683803.199869424

Cost for cross validation dataset:
0.00490850113873416

Cost for test dataset:
0.004976562444356282

Used R^2 score as an evaluation metric to check for testing the performance of the model.

R^2 Score for training dataset: 0.99999999922728

R^2 Score for cross validation dataset: 0.99999999924968


R^2 Score for test dataset: 0.999999999221306

Polynomial Regression

For polynomial regression, the first decision to make was the degree of the polynomial to be used. I considered the powers of the initial three features as (u, v, w) and used list comprehension to solve for $u + v + w = i$ where $i = (1, 2, \dots, n)$, n being the degree of polynomial taken as input. So, in this way, I engineered new features using the given features. Then, I shuffled the dataset, did a 0.8 : 0.1 : 0.1 split, and Z-score normalized the features.

Keeping in mind the increased complexity of the dataset, I implemented L2 regularization. Then, I ran gradient descent for different values of degree n , learning rate and regularization parameter (λ), finally I found optimal results for $n=6$, learning rate of 0.25 and 15000 iterations. I also tried different values for λ and observed the cost on cross validation set and realized that the model wasn't suffering high variance issue, so I decided to set λ to 0. The cost (Mean Squared Error) saturated out to $1.607e-16$.

Cost vs Iterations:



Initial Cost: $7.230e+13$

Cost for cross validation dataset:
 $1.0080838554572531e-16$

Cost for test dataset:
 $1.0732292957743687e-16$

Used R^2 score as an evaluation metric to check for testing the performance of the model.

R^2 Score for training dataset: 1.0

R^2 Score for cross validation dataset: 1.0

R^2 Score for test dataset: 1.0

Logistic Regression

I did a 0.8 : 0.1 : 0.1 (train : cross validation : test) split for Logistic Regression and then scaled the data (pixel values) to between 0 and 1 by dividing them by 255 (maximum value for a pixel). Then, I one hot encoded the labels in each set and stored them in different variables.

Further, since the dataset had multiple classes (10), the optimal choice would have been to use the softmax function, but since the task was to build a model for 'logistic regression', I felt it implied using the sigmoid function, thus I decided to go ahead with the 'One vs rest' approach. I considered trying the 'One vs One' approach but eventually decided that training 10C2 models would take a lot of time, so dropped that idea.


So, I vectorized the 10 models for the One vs rest approach by taking the dimensions of the weight as (number of features, number of classes).

I implemented L2 regularization but found that high variance was not an issue in the model so set the lambda to 0.

I ran gradient descent on different values of learning rate and got the following results:

A cost (Binary Cross Entropy) of around 0.02416 with a learning rate of 5 and 5000 iterations

Cost vs Iterations:



Initial Cost: 0.6931471805599451

Accuracy for training dataset:
0.9745833333333334

Accuracy for training dataset:
0.9716666666666667

Neural Network


I shuffled the dataset and did a 0.8:0.1:0.1 (train: cross validation: test) split and then scaled the pixel values to between 0 and 1 by dividing the dataset by 255 (maximum value for a pixel), one hot encoded the labels and stored them separately. Then I initialized the weights with the 'He initialization', to avoid the issue of vanishing and exploding gradients. I saved the parameters in a dictionary for easy access ahead.

So, I used a model with three hidden ReLU activation layers having 512, 256 and 128 neurons and used the softmax function in the output layer having 10 neurons corresponding to 10 classes.

Then I ran gradient descent but it turned out to be pretty slow, so I decided to implement mini batch gradient descent which was much faster and even better in terms of adding variance with each mini batch. I also implemented L2 regularization but later figured out that there was no high variance issue in the model so I set the lambda to 0. I used a mini batch size of 64 examples.

The cost (Sparse Categorical Cross Entropy) came around $1.6e-05$ after 50 epochs, with a learning rate of 0.25.

Cost vs epochs:



Initial Cost: 4.757739175084508

Accuracies:

Accuracy for training dataset:
1.0

Accuracy for cross validation
dataset:
0.9886666666666667

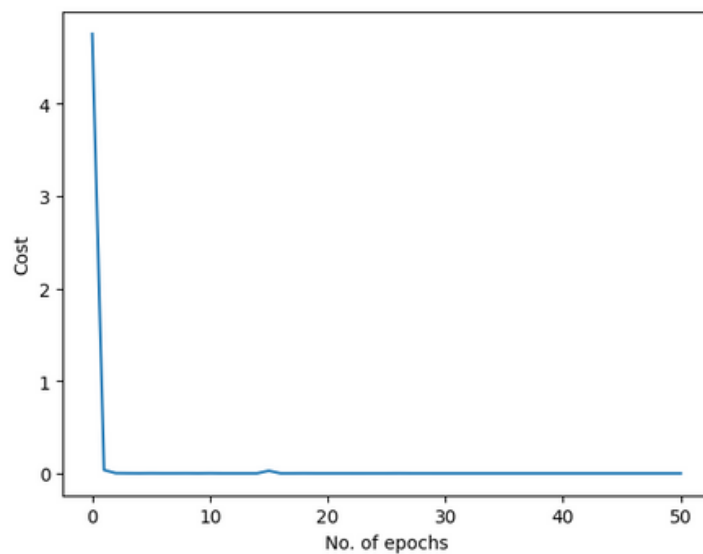
Accuracy for test dataset:
0.989

Then, out of curiosity, I tried optimizers like gradient descent with momentum, Adam (Adaptive moment estimation); which helped in improving the accuracies.

Adam:

Got optimal results with learning rate of 0.001 and 50 epochs, beta1 of 0.9 and beta2=0.999 with a cost of around $4.6e-09$.

Cost vs epochs:



Initial Cost: 4.757739175084508

Accuracies:

Accuracy for training dataset:

1.0

Accuracy for cross validation dataset:

0.9916666666666667

Accuracy for test dataset:


0.988

Gradient Descent with momentum:

Tried different values of learning and epochs and got the following results-

A cost of around $5.86e-06$ with a learning rate of 0.5, beta of 0.9 and 30 epochs

Cost vs epochs:



Initial Cost: 4.757739175084508

Accuracies:

Accuracy for training dataset:

1.0

Accuracy for cross validation dataset:


0.992

Accuracy for test dataset:

0.989

A cost of around $1.9\text{e-}06$ with a learning rate of 0.5, beta of 0.9 and 100 epochs.

Cost vs epochs:



Initial Cost: 4.757739175084508

Accuracies:

Accuracy for training dataset:

1.0

Accuracy for cross validation dataset:

0.9923333333333333

Accuracy for test dataset:

0.9893333333333333

While using Adam and momentum, the cost doesn't converge to the global minima but rather revolves around it, thus I tried using decay for reducing the learning rate when the cost nears the minima, so that it converges. But I figured out that the results were better or almost similar without the decay, so I set the decay rate to 0.

It is to be noted that the Adam optimizer gives much lesser cost but a lower accuracy than momentum.

KNN


I shuffled the data and split it into the train set and the test set in a 0.8:0.2 ratio, then I scaled the pixel values to between 0 and 1 by dividing the data by 255 (maximum value of a pixel).

First I tried the 'one loop method', in which calculating the distance of a test data point from each point in the train data was vectorized but had to run a loop for all the data points in test set. This method was taking a lot of time, about 5-10 minutes.

Then, I searched and studied the 'no loop method' which vectorizes the process of calculating the distances of all data points in test set from the train set.

After calculating the distances, I used `np.argsort` to get the indices of K shortest distances for all test data points, and stored the labels corresponding to these nearest indices. Then I calculated and assigned the most common label in the K labels for each test data point.

Accuracy vs K:



So, the best accuracy is achieved for the value of 1 for the hyperparameter K, and it is observed that each successive maxima is achieved at odd values of K.


Accuracy at K=1: 0.9858333333333333

K-Means

I first Z-score normalized the dataset. Then I randomly initialized the clusters by choosing any K data points from the dataset.

Then I iterated over the steps of finding the closest centroids for each data point and then updating each centroid with the mean of the data points assigned to it.

Cost vs K:



Using the elbow method, it can be identified that K=2 and 6 are forming elbows, but K=6 should be the optimal choice considering that it is more practical.